

CS-E4820 Machine Learning: Advanced Probabilistic Methods (spring 2021)

Pekka Marttinen, Santosh Hiremath, Tianyu Cui, Yogesh Kumar, Zheyang Shen, Alexander Aushhev, Khaoula El Mekkaoui, Shaoxiong Ji, Alexander Nikitin, Sebastiaan De Peuter, Joakim Järvinen.

Exercise 4, due on Tuesday February 16th at 23:00.

Problem 1: ML-II for a linear model

Fit the Bayesian linear parameter model to a given data 'ex4_1_data.txt' using the ML II approach. Optimize the hyperparameters α and β using grid search. Complete the template given below with your own code. Make predictions for the test data using the fitted model and compute the mean squared error for test data. Also plot the data and the fitted model.

```
[1]: # Starter code for problems 1 and 2.
# NOTE: variables defined in this cell are used in code templates for
→problem 1 and 2.

import numpy as np
import matplotlib.pyplot as plt

# INITIALIZATION
# Load the observations
data = np.loadtxt('/coursedata/ex4_1_data.txt')
x_obs = data[:,0]
y_obs = data[:,1]

# Training data
N_train = 40
x_train = x_obs[:N_train]
y_train = y_obs[:N_train]

# Validation data
N_valid = 10
x_valid = x_obs[N_train:N_train+N_valid]
y_valid = y_obs[N_train:N_train+N_valid]

# Testing data
N_test = 10
x_test = x_obs[N_train+N_valid:]
y_test = y_obs[N_train+N_valid:]

x_range = (-5, 5) # Possible values of x are in this range

# Basis function parameters
num_basis_functions = 11
centers = np.linspace(x_range[0], x_range[1], num_basis_functions)
```

```

    lambdaval = 0.17
    # You can use here assume the correct basis function centers and lambda .
→...
def rbf(x, centers, lambdaval):
    # Radial Basis Function output for input x
    #
    # Inputs:
    # x : input points (one-dimensional array)
    # centers : basis function centers (one-dimensional array)
    # lambdaval : basis function width (scalar)
    #
    # Output:
    # Radial Basis Functions evaluated at x (two-dimensional array with
→len(x)
    #
    #                                     rows and len(centers) columns)
    d = x[:,np.newaxis] - centers[np.newaxis,:]
    y = np.exp(-0.5 * (d ** 2) / lambdaval)
    return y

```

```

[2]:
    # Template for problem 1
def bayesian_linear_regression(phi_x, y, alpha, beta):
    # Bayesian linear parameter model
    #
    # Inputs:
    # phi_x : the basis function applied to x-data (two-dimensional array)
    # y : y-data (one-dimensional array)
    # alpha : the precision of the weight prior distribution (scalar)
    # beta : the precision of the assumed gaussian noise (scalar)
    #
    # Output:
    # the posterior mean, the posterior covariance, the log marginal
→likelihood

```

```

    N, B = phi_x.shape

```

```

    # Add here code to compute:
    # m = ? # EXERCISE: the posterior mean of w
    # S = ? # EXERCISE: the posterior covariance of w
    # S_inv = ? # EXERCISE: the inverse of S

```

```

    ### BEGIN SOLUTION

```

```

    S_inv = alpha * np.eye(B) + beta * np.dot(phi_x.T, phi_x)
    S = np.linalg.inv(S_inv)
    d = beta * np.dot(phi_x.T, y)
    m = np.dot(S, d)

```

```

    ### END SOLUTION

    # Note: This is a corrected version of equation 18.1.19 from Barbers book
    d = beta * np.dot(phi_x.T, y)
    log_likelihood = 0.5 * (-beta * np.dot(y, y) + d @ S @ d + np.log(np.
→linalg.det(2 * np.pi * S)) +
    B * np.log(alpha) + N * np.log(beta) - N * np.log(2 * np.pi))
    return m, S, log_likelihood

    # Specify possible values for the alpha and beta parameters to test
    alphas = np.logspace(-3, 3, 100)
    betas = np.logspace(-3, 3, 100)

    # Grid search over possible values of alpha and beta
    best_log_likelihood = -np.inf # optimal parameter values maximize the
→log likelihood
    for alpha in alphas:
    for beta in betas:
        # Use here functions rbf and bayesian_linear_regression to compute the
        # log marginal likelihood for given alpha and beta

        # What are the optimal values of alpha and beta, that maximize the
→marginal likelihood?

        # best_alpha = ? # EXERCISE
        # best_beta = ? # EXERCISE

        # Fit the model one more time using the optimal alpha and beta and the
→training data
        # to get m for the optimal model

        # best_m = ? # EXERCISE

    ### BEGIN SOLUTION
    phi = rbf(x_train, centers, lambdaval)
    m, S, log_likelihood = bayesian_linear_regression(phi, y_train, alpha,
→beta)
    if log_likelihood > best_log_likelihood:
        best_alpha = alpha
        best_beta = beta
        best_m = m
        best_log_likelihood = log_likelihood
    ### END SOLUTION

    # Compute the final regression function
    x_coord = np.linspace(x_range[0], x_range[1], 100)

```

```

# Compute the predicted values for inputs in x_coord using best_m
# y_mean = ? # EXERCISE

### BEGIN SOLUTION
y_mean = np.dot(rbf(x_coord, centers, lambdaval), best_m)
### END SOLUTION

# Plot the final learned regression function, together with the samples
plt.plot(x_coord, y_mean, label="learned model")
plt.plot(x_train, y_train, 'kx', label="training data")
plt.plot(x_test, y_test, 'rx', label="testing data")

# Make predictions for inputs in the test data, so that you get
# predictions 'y_pred' for inputs in x_test.
# y_pred = ? #EXERCISE

### BEGIN SOLUTION
y_pred = np.dot(rbf(x_test, centers, lambdaval), best_m)
### END SOLUTION

# print(y_pred)


# Plot the predictions
plt.plot(x_test, y_pred, 'gx', label="testing predictions")

# Compute the mean squared prediction error for the test data.
# mse_test = ???

### BEGIN SOLUTION
mse_test = np.mean((y_pred - y_test) ** 2)
### END SOLUTION

plt.legend()
plt.title("ML-II:  $\alpha$ =%.3f,  $\beta$ =%.3f, mse=%.4f" %
(best_alpha, best_beta, mse_test))
plt.show()

```



Problem 2: Optimizing hyperparameters with validation set

As in problem 1, fit the Bayesian linear parameter model to a given data 'ex4_1_data.txt', but optimize the hyperparameters α and β by dividing the training data into training and validation sets, and selecting the values of α and β that minimize the mean squared error for the validation set. Make predictions for the test data using the fitted model and compute the mean squared error for test data. Plot the data and the fitted model.

```
[3]: # Template for problem 2
# Specify possible values for the alpha and beta parameters to test
alphas = np.logspace(-3, 3, 100)
betas = np.logspace(-3, 3, 100)

# Grid search over possible values of alpha and beta
mse_valid = np.zeros((len(alphas), len(betas)))
for a, alpha in enumerate(alphas):
    for b, beta in enumerate(betas):

# Use here functions rbf and bayesian_linear_regression to fit the
# model and compute the prediction error (using mean squared error)
# for the validation data

### BEGIN SOLUTION
phi = rbf(x_train, centers, lambdaval)
```

```

    m, S, log_likelihood = bayesian_linear_regression(phi, y_train, alpha,
→beta)
    y_pred_valid = np.dot(rbf(x_valid, centers, lambdaval), m)
    mse_valid[a, b] = np.mean((y_pred_valid - y_valid) ** 2)
    ### END SOLUTION

    # What are the optimal values of alpha and beta, that minimize the
→prediction error in the validation data?
    # best_alpha = ? #EXERCISE
    # best_beta = ? #EXERCISE

    ### BEGIN SOLUTION
    (best_a, best_b) = np.unravel_index(np.argmin(mse_valid), mse_valid.
→shape)
    best_alpha = alphas[best_a]
    best_beta = betas[best_b]
    ### END SOLUTION

    # Fit the model one more time using the optimal alpha and beta and all
→data
    # available for model fitting (both training and validation sets)
    # best_m = ? #EXERCISE

    ### BEGIN SOLUTION
    best_m, _, _ = bayesian_linear_regression(phi, y_train, best_alpha,
→best_beta)
    ### END SOLUTION

    x_coord = np.linspace(x_range[0], x_range[1], 100)

    # Compute the predicted values for inputs in x_coord using best_m
    # y_mean = ? #EXERCISE

    ### BEGIN SOLUTION
    y_mean = np.dot(rbf(x_coord, centers, lambdaval), best_m)
    ### END SOLUTION

    # Plot the final learned regression function, together with the samples
    plt.plot(x_coord, y_mean, label="learned model")
    plt.plot(x_train, y_train, 'kx', label="training data")
    plt.plot(x_valid, y_valid, 'bx', label="validation data")
    plt.plot(x_test, y_test, 'rx', label="testing data")

    # Make predictions for inputs in the test data, so that you get
    # predictions 'y_pred' for inputs in x_test.
    # y_pred = ? ? # EXERCISE

```

```

### BEGIN SOLUTION
y_pred = np.dot(rbf(x_test, centers, lambdaval), best_m)
### END SOLUTION

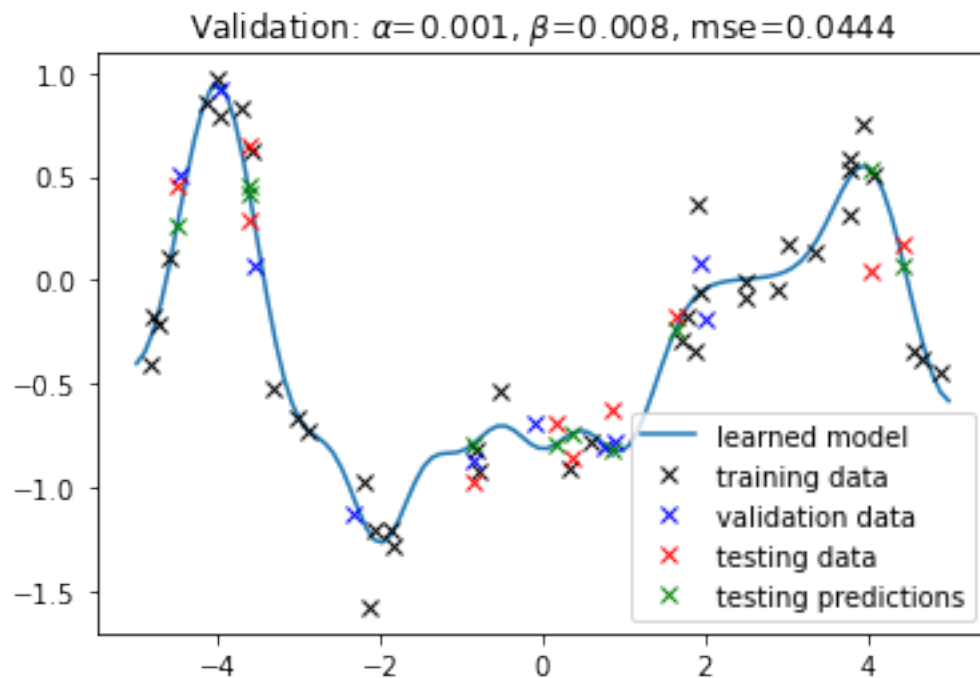
# Plot the predictions
plt.plot(x_test, y_pred, 'gx', label="testing predictions")

# Compute the mean squared prediction error for the test data.
# mse_test = ? # EXERCISE

### BEGIN SOLUTION
mse_test = np.mean((y_pred - y_test) ** 2)
### END SOLUTION

plt.legend()
plt.title("Validation:  $\alpha=0.001$ ,  $\beta=0.008$ , mse=0.0444" %
(best_alpha, best_beta, mse_test))
plt.show()

```



Problem 3: Poisson regression with Laplace approximation

Poisson regression can be used to model count data. A Poisson regression model can be defined as

$$y_i \mid \boldsymbol{\theta} \sim \text{Poisson}(\exp(\boldsymbol{\theta}^T \mathbf{x}_i)) \quad (1)$$

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I}) \quad (2)$$

where y_i are the observed counts, \mathbf{x}_i the related covariates, $i = 1, \dots, N$, and $\boldsymbol{\theta}^T$ are the regression weights. In this exercise, we approximate the posterior $p(\boldsymbol{\theta} \mid \mathbf{y})$ using the Laplace approximation. We will do this in two steps. In the first step we will derive the gradient $-\nabla \log p(\boldsymbol{\theta} \mid \mathbf{y})$ and in the second step we will write Laplace approximation. And finally, we will look compare the true density with the laplace approximation.

(a) Derive the gradient $-\nabla \log p(\boldsymbol{\theta} \mid \mathbf{y})$ and the Hessian $\mathbf{H} = -\nabla \nabla \log p(\boldsymbol{\theta} \mid \mathbf{y})$ needed for the Laplace approximation.

(b) Write the Laplace approximation as the density of a Gaussian distribution. What is the mean and the covariance matrix of this distribution?

(c) Compare the Laplace approximation to the true posterior (computed using numerical integration), in a case where we have one-dimensional covariates only. Use data given in the file 'ex4_4_data.txt' and hyperparameter value $\alpha = 10^{-2}$. Plot the two posteriors and the true value $\theta = \pi/4$ used to generate the data. Also plot the data with the regression line $\hat{y}_i = \exp(\hat{\theta} x_i)$ using the MAP estimate $\hat{\boldsymbol{\theta}}$. The code template below at the end of the notebook will help with this.

Solution

(a) The negative log posterior is given by

$$E(\boldsymbol{\theta}) = -\log p(\boldsymbol{\theta} \mid \mathbf{y}) = -\sum_i \left(y_i \cdot \boldsymbol{\theta}^T \mathbf{x}_i - \exp(\boldsymbol{\theta}^T \mathbf{x}_i) - \log y_i! \right) + \frac{\alpha}{2} \boldsymbol{\theta}^T \boldsymbol{\theta}.$$

Note that in the current context, E stands for *energy* (a term which derives from physics), **not** expectation. The gradient with respect to the parameters $\boldsymbol{\theta}$ is then

$$\frac{\partial E}{\partial \boldsymbol{\theta}} = -\sum_i \left(y_i \mathbf{x}_i - \mathbf{x}_i \exp(\boldsymbol{\theta}^T \mathbf{x}_i) \right) + \alpha \boldsymbol{\theta}$$

and the Hessian

$$\mathbf{H} = \frac{\partial^2 E}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}^T} = +\sum_i \mathbf{x}_i \mathbf{x}_i^T \exp(\boldsymbol{\theta}^T \mathbf{x}_i) + \alpha \mathbf{I}.$$

(b) The Laplace approximation is a Gaussian density $\mathcal{N}(\boldsymbol{\theta} \mid \boldsymbol{\theta}^*, \mathbf{S})$, where the mean $\boldsymbol{\theta}^*$ is the value of the parameter for which the gradient of the log posterior is 0, and the covariance matrix $\mathbf{S} = \mathbf{H}^{-1} \mid_{\boldsymbol{\theta}^*}$ is the inverse of the Hessian evaluated at $\boldsymbol{\theta}^*$.

(c) Run the 3(c) solution code to see the plots.

[4]: `# Template for problem 3(c)`


```

# NOTE: starter code at the begining of the notebook and function
→ 'bayesian_linear_regression'
#         is required to run this cell

import numpy as np
import matplotlib.pyplot as plt
from scipy.special import factorial
from scipy.stats import norm

# get some data
data = np.loadtxt('/coursedata/ex4_4_data.txt')
x = data[:,0]
y = data[:,1]

theta_true = np.pi / 4 # true parameter used to generate the data
alpha = 1e-2 # prior's parameter

# compute Laplace approximation
theta_lapl = 0.5 # initial
# iterate to optimum with newton's method to find the MAP estimate for
→ theta
for iter in range(100):
    # compute gradient
    # grad = ???

    ### BEGIN SOLUTION
    grad = - np.dot(x, y) + np.dot(x, np.exp(x * theta_lapl)) + alpha *
→ theta_lapl
    ### END SOLUTION

    # compute Hessian
    # H = ???

    ### BEGIN SOLUTION
    H = np.dot(x, x * np.exp(x * theta_lapl)) + alpha
    ### END SOLUTION

    theta_lapl = theta_lapl - grad / H # do newton step

# compute Hessian at optimum
→ # H = ???

    ### BEGIN SOLUTION
    H = np.dot(x, x * np.exp(x * theta_lapl)) + alpha
    ### END SOLUTION
print(H)

```

```

difference = theta_lapl - theta_true

# plot posterior densities
theta = np.linspace(0.55, 0.95, 1000)
post_true = np.zeros(len(theta))
for i in range(len(theta)):
    # log posterior:
    post_true[i] = (np.dot(y, x * theta[i]) - np.sum(np.exp(x * theta[i]) -
np.log(factorial(y))) - 0.5*alpha*np.dot(theta[i], theta[i]))

M = np.amax(post_true)
post_true = np.exp(post_true-M) / np.sum(np.exp(post_true-M)) /
→(theta[1]-theta[0]) # normalize

# compute approximative density at the points 'theta'
# Hint: you can use norm.pdf from scipy.stats
# post_laplace = ???

### BEGIN SOLUTION
post_laplace = norm.pdf(theta,theta_lapl,np.sqrt(1/H))
### END SOLUTION

plt.figure(1)
plt.plot(theta, post_true, '-k', label="True posterior")
plt.plot(theta, post_laplace, '-.r', label="Laplace approximation")
plt.plot(theta_true, 0, 'o', label="True value")
plt.xlim(0.55, 0.95)
plt.xlabel('$\\theta$')
plt.title('Posterior $p(\\theta|y)$')
plt.legend()

plt.figure(2)
plt.plot(x, y, 'o', x, np.exp(theta_lapl*x), '-r')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Data')
plt.show()

```

